



THE UNIVERSITY of EDINBURGH
informatics

Coursework 1

Operating Systems
Spring 25

Overview

Introduction

- In this coursework you shall explore the data structures and mechanisms that the Linux kernel uses to manage processes, threads, and scheduling decisions.



Important Dates

- Coursework will be released on 11 Feb
- Coursework due on 25 Feb 2025 12:00

Scoring

- This coursework contributes to 17% of your overall grade.
- There are 3 tasks, scored independently.

Component	Weight
Task 1	20
Task 2	30
Task 3	50
Maximum Score	100

Instructions

Base Kernel

- We provide you a base kernel.
- Download it from [here](#).
- You must work on this kernel.
- This kernel contains incomplete function definitions which you must complete.
- The following CONFIG options must be set
 - CONFIG_SMP
 - CONFIG_SCHED_STATS
 - CONFIG_SCHED_DEBUG


Submission

- You must restrict all your work to a single file:
kernel/sched/core.c
- Submit your core.c in a zip archive named
CW1_<your student UUN>.zip
- Please DO NOT MODIFY ANY OTHER FILE.

Specifications

Task 1 – What's My Ancestry?

- Implement a New Syscall (ID: 463)
- Syscall(463, pid, n)
 - Pid: Target Process
 - N: Max Height
- Should Return the pid of the nth ancestor
- Find Incomplete definition (ancestor_pid) in core.c



```
sys_ancestor_pid(p5, 0) = p5.  
sys_ancestor_pid(p5, 1) = p4.  
sys_ancestor_pid(p5, 2) = p3.  
sys_ancestor_pid(p4, 2) = -ESRCH.  
sys_ancestor_pid(p4, 3) = -ESRCH.
```

Task 1 – Testing

**NOTE: if pid argument is 0,
Use PID of the caller process**

```
#include <unistd.h>
#include <sys/syscall.h>
#include <stdio.h>
#include <sys/wait.h>

#define SYSCALL1 463 // Use the assigned syscall number

int main() {
    pid_t parent = getpid(), gparent = getppid();
    pid_t argpid = 0;

    pid_t pid = fork(); // Fork a new process

    if (pid == -1) {
        perror("fork");
        return 1;
    }

    if (pid == 0) {
        // Call the syscall from the child process
        for (unsigned int n = 0; n < 10; ++n)
            printf("CHILD: %d->%ld\n", n, syscall(SYSCALL1, argpid, n));
        return 0;
    } else {
        // Parent process
        wait(NULL); // Wait for the child to finish
        printf("PARENT: chld:%d me(par):%d gp:%d\n", pid, parent, gparent);
    }

    return 0;
}
```

CHILD: 0->1100

CHILD: 1->1099

CHILD: 2->977

CHILD: 3->822

CHILD: 4->1

CHILD: 5->0

CHILD: 6->-1

CHILD: 7->-1

CHILD: 8->-1

CHILD: 9->-1

PARENT: chld:1100 me(par):1099 gp:977

Task 1 – Where to start?

- Find the incomplete syscall inside core.c
- COMPLETE IT!
- Check out:
 - struct task_struct in *include/linux/sched.h*
 - Find pointer to parent task inside task_struct

```
8135  /**
8136   * CW1
8137   * sys_ancestor_pid – get PID of the nth ancestor of process
8138   * @pid: pid of the process
8139   * @n: nth ancestor
8140   *
8141   * Return: PID on success. Error otherwise.
8142   */
8143  SYSCALL_DEFINE2(ancestor_pid, pid_t, pid, unsigned int, n)
8144  {
8145      return 0;
8146  }
8147
```


Task 2 – TrickleDown Niceness

- Implement a syscall ID 464
- Syscall(464, n)
 - Increment caller's nice value by n
 - Increment its children's niceness by $n/2$
 - Increment its grand-children by $n/4$
 - Increment its great-grand-children by $n/8$
 - So on

Task 2

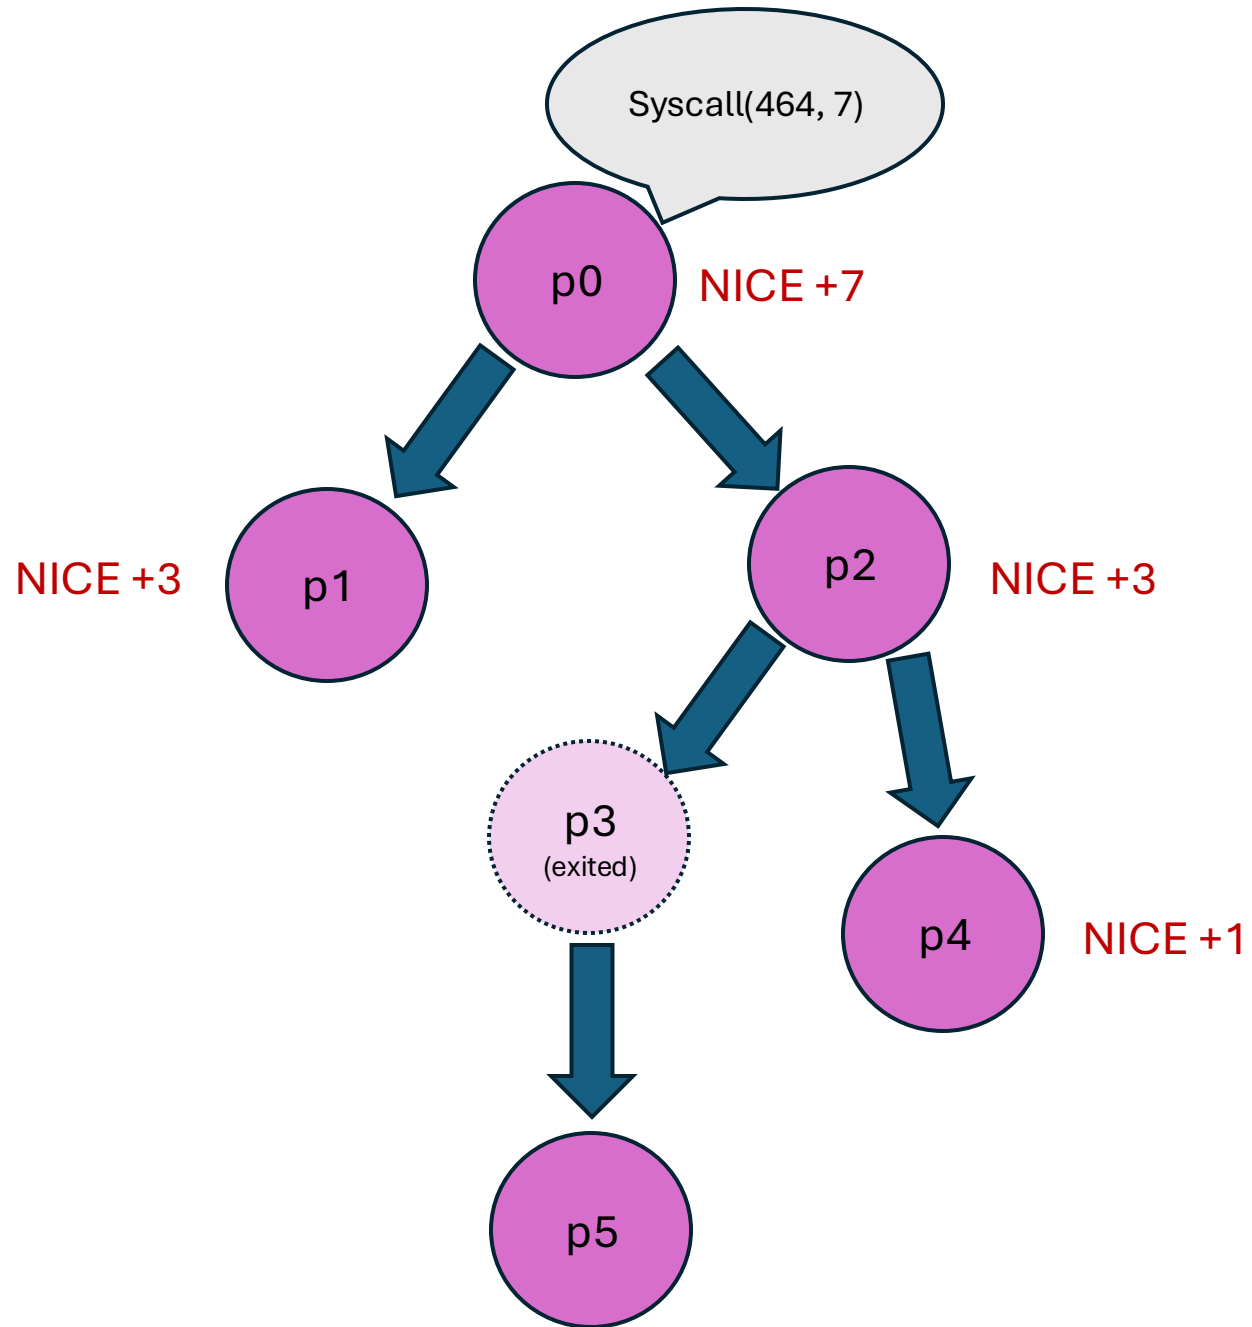
Example

Initially all processes
are at NICE=0




Task 2

Example




Task 2

Example



Task 2

Example



Task 2 - Testing

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <sys/types.h>
#include <sys/resource.h>

#define SYSCALL2 464

int main()
{
    pid_t pid1 = -1, pid2 = -1;


    pid1 = fork();
    if (pid1 == 0) {
        pid2 = fork();
    }

    printf("pid1:%d, pid2:%d, ni:%d\n",
        pid1, pid2, getpriority(PRIO_PROCESS, 0));

    sleep(1);
    if (pid2 == -1)
        syscall(SYSCALL2, 5);
    sleep(1);

    printf("pid1:%d, pid2:%d, ni:%d\n",
        pid1, pid2, getpriority(PRIO_PROCESS, 0));

    return 0;
}
```



OUTPUT

pid1:1273, pid2:-1, ni:0
pid1:0, pid2:1274, ni:0
pid1:0, pid2:0, ni:0

Before
syscall

pid1:1273, pid2:-1, ni:5
pid1:0, pid2:1274, ni:2
pid1:0, pid2:0, ni:1

After
syscall

Task 2 – Where to start?

- Find the incomplete syscall inside core.c
- COMPLETE IT!
- Check out:
 - struct task_struct in *include/linux/sched.h*
 - Find pointer to children tasks inside task_struct

```
7358  /**
7359  * CW1
7360  * sys_propagate_nice - trickle-down nice-increment to descendants
7361  * @increment: nice-increment for calling process
7362  *
7363  * Return: 0 on success. Error otherwise.
7364  */
7365  SYSCALL_DEFINE1(propagate_nice, int, increment)
7366  {
7367      return 0;
7368  }
```

Task 3 – Scheduling History

- Print scheduler stats inside `/proc/<PID>/schedstat`
 - By default it prints 3 numbers
 - In your base kernel, it will have a fourth string (enclosed by '[]')
 - Implement this functionality: This fourth string should print the list of CPUs this process was scheduled on in this epoch (epoch=10 seconds of runtime)
 - Time spent waiting on a runqueue should not be included in epoch! Only RUNTIME is counted.

Task 3 - Where to Start?

- Struct `task_struct` has 2 new members
 - `used_cpus` is a `cpumask_t`
 - When a process is scheduled on a CPU, set the CPU in `used_cpus`
 - `Epoch_ticks` is the number of scheduler ticks this process endured in this epoch.
 - Use this to track the end of an epoch (`used_cpus` must be reset every epoch).
- `Core.c` contains a macro `TICKS_PER_EPOCH`
 - Equals Number of scheduler ticks in 10 secs (epoch)

```
788      /* Tracking a recently used CPU at
789      * used CPU that may be idle.
790      */
791      int      recent_used_cpu;
792      int      wake_cpu;
793
794      /* CW1 */
795      cpumask_t used_cpus;
796      unsigned int epoch_ticks;
797  #endif
798      int      on_rq;
799
800      int      prio;
```

```
117      EXPORT_TRACEPOINT_SYMBOL_GPL(sched_compute_energy_
118
119      DEFINE_PER_CPU_SHARED_ALIGNED(struct rq, runqueues
120
121      /* CW1 - an epoch equals 10 seconds */
122      #define TICKS_PER_EPOCH (10 * HZ)
123
124      #ifdef CONFIG_SCHED_DEBUG
125      /*
126      * Debugging: various feature bits
127      *
```

Best of Luck!